

# Intro to Deep Learning

## Big Data y Machine Learning para Economía Aplicada

Ignacio Sarmiento-Barbieri

Universidad de los Andes

# Agenda

- 1 Single Layer Neural Networks
- 2 Activation Functions
- 3 Output Functions
- 4 Training the network

# Deep Learning: Intro


- ▶ Linear Models may miss the nonlinearities that best approximate  $f^*(x)$
- ▶ Neural networks are simple models.
- ▶ The model has **linear combinations** of inputs that are passed through **nonlinear activation functions** called nodes (or, in reference to the human brain, neurons).

# Deep Learning: Intro

- ▶ Let's start with a familiar and simple model, the linear model

# Deep Learning: Intro

- ▶ Let's start with a familiar and simple model, the linear model



# Single Layer Neural Networks

- ▶ A neural network takes an input vector of  $p$  variables

$$X = (X_1, X_2, \dots, X_p) \quad (1)$$

- ▶ and builds a nonlinear function  $f(X)$  to predict the response  $y$ .


$$y = f(X) + u \quad (2)$$

- ▶ The Single layer NN model has the form

$$f(X) = f^{(output)}(g(X)) \quad (3)$$

- ▶ where  $g$  is the activation function in the hidden unit
- ▶ the second layer,  $f^{(output)}$  is the output layer of the network

# Single Layer Neural Networks



# Single Layer Neural Networks

- ▶ NN are made of **linear combinations** of inputs that are passed through **nonlinear activation functions**




# Worked Example I: Single Layer Neural Networks

- ▶ 2 Predictors:  $p = 2$ ,  $X = (X_1, X_2)$
- ▶ 2 Nodes:  $K = 2$ ,  $A_1(X)$  and  $A_2(X)$
- ▶ Non-linear activation function  $g(z) = z^2$
- ▶ Want to predict a number  $\in \mathbb{R}$ : identity output function ( $f^{(output)} : \mathbb{R} \rightarrow \mathbb{R}$  such that  $f^{(output)}(x) = x$ )

# Worked Example I: Single Layer Neural Networks

- ▶ 2 Predictors:  $p = 2$ ,  $X = (X_1, X_2)$
- ▶ 2 Nodes:  $K = 2$ ,  $A_1(X)$  and  $A_2(X)$
- ▶ Non-linear activation function  $g(z) = z^2$
- ▶ Want to predict a number  $\in \mathbb{R}$ : identity output function ( $f^{(output)} : \mathbb{R} \rightarrow \mathbb{R}$  such that  $f^{(output)}(x) = x$ )



# Worked Example I: Single Layer Neural Networks

- ▶ 2 Predictors:  $p = 2$ ,  $X = (X_1, X_2)$
- ▶ 2 Nodes:  $K = 2$ ,  $A_1(X)$  and  $A_2(X)$
- ▶ Non-linear activation function  $g(z) = z^2$
- ▶ Want to predict a number  $\in \mathbb{R}$ : identity output function ( $f^{(output)} : \mathbb{R} \rightarrow \mathbb{R}$  such that  $f^{(output)}(x) = x$ )

$$f(X) = f^{(output)}(g(X)) \quad (4)$$

# Worked Example I: Single Layer Neural Networks

- ▶ 2 Predictors:  $p = 2$ ,  $X = (X_1, X_2)$
- ▶ 2 Nodes:  $K = 2$ ,  $A_1(X)$  and  $A_2(X)$
- ▶ Non-linear activation function  $g(z) = z^2$
- ▶ Want to predict a number  $\in \mathbb{R}$ : identity output function ( $f^{(output)} : \mathbb{R} \rightarrow \mathbb{R}$  such that  $f^{(output)}(x) = x$ )

$$f(X) = f^{(output)}(g(X)) \quad (4)$$

$$f(X) = \beta_0 + \sum_{k=1}^2 \beta_k \left( w_{k0} + \sum_{j=1}^2 w_{kj} X_j \right)^2 \quad (5)$$

# Why not linear activation functions?

## Worked Example II : The "Exclusive OR (XOR)" Function

- ▶ The exclusive disjunction of a pair of propositions,  $(p, q)$ , is supposed to mean that  $p$  is true or  $q$  is true, but not both
- ▶ It's truth table is:

q	p	y
0	0	0
0	1	1
1	0	1
1	1	0

# Agenda

- 1 Single Layer Neural Networks
- 2 Activation Functions**
- 3 Output Functions
- 4 Training the network

# Activation Functions

## Sigmoid Function (Logit)


- ▶ The sigmoid function transforms its inputs, for which values lie in the domain  $\mathbb{R}$ , to outputs that lie on the interval  $(0, 1)$ .
- ▶ For that reason, the sigmoid is often called a squashing function: it squashes any input in the range  $(-\infty, \infty)$  to some value in the range  $(0, 1)$ :

$$\text{sigmoid}(x) = \frac{1}{1 + \exp(-x)}.$$



# Activation Functions

## Sigmoid Function (Logit)



- When attention shifted to gradient based learning, the sigmoid function was a natural choice because it is smooth and differentiable.

# Activation Functions

## ReLU Function


### ► ReLU Function

- The most popular choice, due to both simplicity of implementation and its good performance on a variety of predictive tasks, is the rectified linear unit (ReLU).
- ReLU provides a very simple nonlinear transformation. Given an element  $x$ , the function is defined as the maximum of that element and 0:

$$\text{ReLU}(x) = \max\{x, 0\}$$

# Activation Functions

- ▶ ReLU function retains only positive elements and discards all negative elements by setting them to 0.
- ▶ It is piecewise linear.



# Neural Networks: Activation Functions

- ▶  $\text{Sigmoid}(x) = \frac{1}{1+\exp(-x)}$
- ▶  $\text{ReLU}(x) = \max\{x, 0\}$
- ▶ Among others ([see more here](#))
- ▶ Hidden unit design remains an active area of research, and many useful hidden unit types remain to be discovered

# Agenda

- 1 Single Layer Neural Networks
- 2 Activation Functions
- 3 Output Functions**
- 4 Training the network

# Output Functions

- ▶ The choice of output unit is related to the problem at hand
  - ▶ Regression
  - ▶ Classification
    - ▶ Binary
    - ▶ Multiclass

# Agenda

- 1 Single Layer Neural Networks
- 2 Activation Functions
- 3 Output Functions
- 4 Training the network

# Training the network

- El objetivo es

$$\hat{f} = \operatorname{argmin}_f \left\{ \sum_{i=1}^n L(y, f(X; \Theta)) \right\} \quad (6)$$


- SNN

$$f(X, \beta, w) = f \left[ \beta_0 + \sum_{k=1}^K \beta_k g \left( w_{k0} + \sum_{j=1}^p w_{kj} X_j \right) \right] \quad (7)$$



# Training the network

Example: House Prices



## ► Equations

### ► Hidden Layer sigmoid (logistic):

$$\text{► } A_1 = \sigma(w_{11} \cdot X_1 + w_{12} \cdot X_2 + w_{10})$$


$$\text{► } A_2 = \sigma(w_{21} \cdot X_1 + w_{22} \cdot X_2 + w_{20})$$

### ► Output Layer, identity output function:


$$\text{► } \hat{y}_i = \beta_0 + \beta_1 \cdot A_1 + \beta_2 \cdot A_2$$

## ► Loss Function $\Rightarrow$ MSE: $\frac{1}{n} \sum_i^n (y_i - \hat{y}_i)^2$

# Training the network



# Training the network



# Training the network

Example: House Prices

